

CS918: LECTURE 6

Vector Representation and Models for Word Embeddings

Arkaitz Zubiaga, 22nd October, 2018

LECTURE 6: CONTENTS

- Vector space models for language representation.
- Word embeddings.
 - SVD: Singular Value Decomposition.
 - Iteration based models.
 - CBOW and skip-gram models.
 - Word2Vec and Glove.

RECAP: STATISTICAL LANGUAGE MODELS

- Goal: compute the **probability of a sequence of words**:

$$P(W) = P(w_1, w_2, w_3, w_4, w_5, \dots, w_n)$$

- Related task: **probability of an upcoming word**:

- $P(w_5 \mid w_1, w_2, w_3, w_4)$

- Both of the above are **language models**.



WARWICK

VECTOR SPACE MODELS

VECTOR REPRESENTATION

- So far we've extracted **n-gram counts**, etc. from texts.
- For **most NLP tasks**, we need a **vector representation**, which can be fed to:
 - Sentiment classifier.
 - Information retrieval system.
 - Question answering system.
 - Etc.

WORDS AS ATOMIC SYMBOLS

- So far, we have viewed **words as (sequences of) atomic symbols**.
 - We have used **edit distance to compute similarity**.
 - **N-grams & LMs** → what may **follow/precede the word?**

WORDS AS ATOMIC SYMBOLS

- So far, we have viewed **words as (sequences of) atomic symbols**.
- This **doesn't tell us anything about semantic similarity**, e.g.:
 - Is “Chinese” closer to “Asian” or to “English”?
 - Are “king” & “queen” more related than “doctor” & “mountain”?

WORDS AS ATOMIC SYMBOLS

- We may identify significant similarity based on word overlap between:
 - “Facebook to fight 'fake news' by asking users to rank trust in media outlets”
 - “Facebook's latest fix for fake news: ask users what they trust”
- Using stemmer/lemmatiser

WORDS AS ATOMIC SYMBOLS

- We may identify significant similarity based on word overlap between:
 - “Facebook to fight 'fake news' by asking users to rank trust in media outlets”
 - “Facebook's latest fix for fake news: ask users what they trust”

————→ Using stemmer/lemmatiser

- But we'll fail when there isn't an overlap:
 - “Zuckerberg announces new feature that crowdsources trustworthiness of news organisations”

NO OVERLAP

WORDS AS ATOMIC SYMBOLS

- Likewise for text classification, e.g.:
 - If classifier learns that:

“Leicester will welcome back Jamie Vardy for their Premier League clash with Watford”
 belongs to the topic “sport”

- We’ll fail to classify the following also as “sport:”

“Blind Cricket World Cup: India beat Pakistan by two wickets in thrilling final to retain title”



NO OVERLAP

WORD VECTORS: ONE-HOT OR BINARY MODEL

- Word represented as: $\{0, 1\}^{|V| \times 1}$ vector, $|V|$ = vocabulary size

e.g. $V = [\text{hotel}, \text{motel}, \text{cat}, \text{dog}]$, $|V| = 4$

hotel = [1, 0, 0, 0]

motel = [0, 1, 0, 0]

cat = [0, 0, 1, 0]

dog = [0, 0, 0, 1]

WORD VECTORS: ONE-HOT OR BINARY MODEL

- Word represented as: $\{0, 1\}^{|V| \times 1}$ vector, $|V|$ = vocabulary size
- Still no notion of similarity, e.g.:

$$(w^{hotel})^T w^{motel} = (w^{hotel})^T w^{cat} = 0$$

BAG-OF-WORDS MODEL

- **Bag-of-words:** $\vec{v} = \{|w_1|, |w_2|, \dots, |w_n|\}$
- Toy example: hello world hello world hello I like chocolate
 $\vec{v} = \{2, 3, 1, 1, 1\}$
- **Widely used**, but largely **being replaced by word embeddings**.
- **Con:** inefficient for large vocabularies.
- **Con:** doesn't capture semantics (each word is an unrelated token)

WORD VECTORS: ONE-HOT OR BINARY MODEL

- **Solution:** why not reduce dimensionality of vector space?

$\mathbb{R}^{N \times 1}$ or (in matrix format) $\mathbb{R}^{N \times |V|}$

from:

| | |
|-------|----------------|
| hotel | = [1, 0, 0, 0] |
| motel | = [0, 1, 0, 0] |
| cat | = [0, 0, 1, 0] |
| dog | = [0, 0, 0, 1] |

to something like:

| | |
|-------|----------|
| hotel | = [1, 0] |
| motel | = [1, 0] |
| cat | = [0, 1] |
| dog | = [0, 1] |

1st dimension = 1 if word is a building

2nd dimension = 1 if word is an animal

now we can relate words!



WORD EMBEDDINGS: SINGULAR VALUE DECOMPOSITION (SVD)

WORD EMBEDDINGS

- Assumptions:
 - We can represent **words as vectors** of some dimension.
 - Each **dimension has some semantic meaning**, unknown a priori, but could be e.g.:
 - Whether it is an object/concept/person.
 - Gender of person.
 - ...

INTUITION OF WORD EMBEDDINGS

- Words with the same context will have similar meaning:

buy a car
buy chocolate
don't **buy**
will you **buy** it?

purchase a car
purchase chocolate
don't **purchase**
will you **purchase** it?

get a car
get chocolate
don't **get**
will you **get** it?

buy, purchase and get occur in equal or very similar contexts
they must have **similar meanings!**

BUILDING A CO-OCCURRENCE MATRIX

- Given as input:
 - A **text/corpus**.
 - An **offset Δ** (e.g. 5 words)
- In a co-occurrence matrix with $|V|$ rows, $|V|$ columns:
 - The $(i, j)^{\text{th}}$ value indicates the **number of times words i and j co-occur within the given offset Δ** .

BUILDING A CO-OCCURRENCE MATRIX

- Examples ($\Delta = 2$ words):
 - We need **to tackle fake news to** keep society informed.
 - How can we build a classifier to **deal with fake news**?
 - **Fake** co-occurs with: **to**(2), **news**(2), **deal**(1), **tackle**(1), **with**(1)
 - **Deal** (with) and **tackle** are different tokens for us.
 - **Frequent occurrence in similar contexts will indicate similarity.**

WORD EMBEDDINGS: WORD-WORD MATRIX

| | i | want | to | eat | chinese | food | lunch | spend |
|---------|----|------|-----|-----|---------|------|-------|-------|
| i | 5 | 827 | 0 | 9 | 0 | 0 | 0 | 2 |
| want | 2 | 0 | 608 | 1 | 6 | 6 | 5 | 1 |
| to | 2 | 0 | 4 | 686 | 2 | 0 | 6 | 211 |
| eat | 0 | 0 | 2 | 0 | 16 | 2 | 42 | 0 |
| chinese | 1 | 0 | 0 | 0 | 0 | 82 | 1 | 0 |
| food | 15 | 0 | 15 | 0 | 1 | 4 | 0 | 0 |
| lunch | 2 | 0 | 0 | 0 | 0 | 1 | 0 | 0 |
| spend | 1 | 0 | 1 | 0 | 0 | 0 | 0 | 0 |

- The table will be **huge (and sparse)** for large $|V|$ (vocabularies).
- We need to **reduce the dimensionality**.

WORD EMBEDDINGS: SVD METHODS

- **SVD: Singular Value Decomposition**
- We **build co-occurrence matrix** ($|V| \times |V|$) with offset Δ .
- We use **SVD to decompose X** as $X = USV^T$, where:
 - $U(|V| \times r)$ and $V(|V| \times r)$ are unitary matrices, and
 - $S(r \times r)$ is a diagonal matrix.
- The **columns of U** (the left singular vectors) are then the **word embeddings of the vocabulary**.

WORD EMBEDDINGS: SVD METHODS

$$\begin{matrix} |V| \\ \left[\begin{array}{c} X \end{array} \right] \end{matrix} = \begin{matrix} |V| \\ \left[\begin{array}{c|c|c} u_1 & u_2 & \dots \end{array} \right] \end{matrix} \begin{matrix} |V| \\ \left[\begin{array}{ccc} \sigma_1 & 0 & \dots \\ 0 & \sigma_2 & \dots \\ \vdots & \vdots & \ddots \end{array} \right] \end{matrix} \begin{matrix} |V| \\ \left[\begin{array}{c} - \\ v_1 \\ - \\ v_2 \\ \vdots \end{array} \right] \end{matrix}$$

Reducing dimensionality by selecting first k singular vectors:

$$\begin{matrix} |V| \\ \left[\begin{array}{c} \hat{X} \end{array} \right] \end{matrix} = \begin{matrix} k \\ \left[\begin{array}{c|c|c} u_1 & u_2 & \dots \end{array} \right] \end{matrix} \begin{matrix} k \\ \left[\begin{array}{ccc} \sigma_1 & 0 & \dots \\ 0 & \sigma_2 & \dots \\ \vdots & \vdots & \ddots \end{array} \right] \end{matrix} \begin{matrix} |V| \\ \left[\begin{array}{c} - \\ v_1 \\ - \\ v_2 \\ \vdots \end{array} \right] \end{matrix}$$



We get $|V|$ vectors of k dimensions each: word embeddings
 e.g. word embedding of word w :

$$\mathbf{WE}(w) = \{\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_k\}$$

We've reduced w 's dimensionality from $|V|$ to k .

SVD EXAMPLE IN PYTHON

Corpus:

I like deep learning. I like NLP. I enjoy flying.

$\Delta = 1$

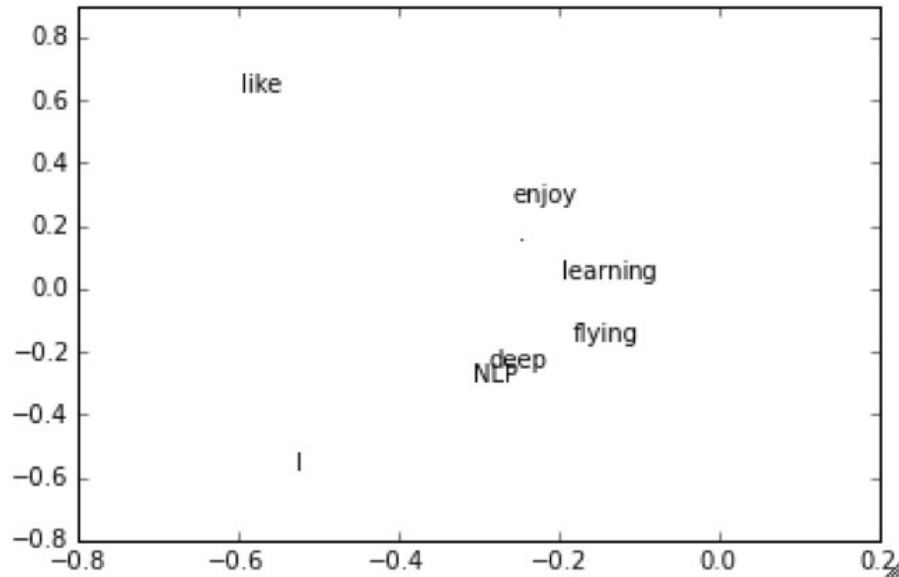
like & I co-occur twice

```
import numpy as np
la = np.linalg
words = ["I", "like", "enjoy",
         "deep", "learning", "NLP", "flying", "."]
X = np.array([[0, 2, 1, 0, 0, 0, 0, 0],
              [2, 0, 0, 1, 0, 1, 0, 0],
              [1, 0, 0, 0, 0, 0, 1, 0],
              [0, 1, 0, 0, 1, 0, 0, 0],
              [0, 0, 0, 1, 0, 0, 0, 1],
              [0, 1, 0, 0, 0, 0, 0, 1],
              [0, 0, 1, 0, 0, 0, 0, 1],
              [0, 0, 0, 0, 1, 1, 1, 0]])

U, s, Vh = la.svd(X, full_matrices=False)
```

PLOTTING SVD EXAMPLE IN PYTHON

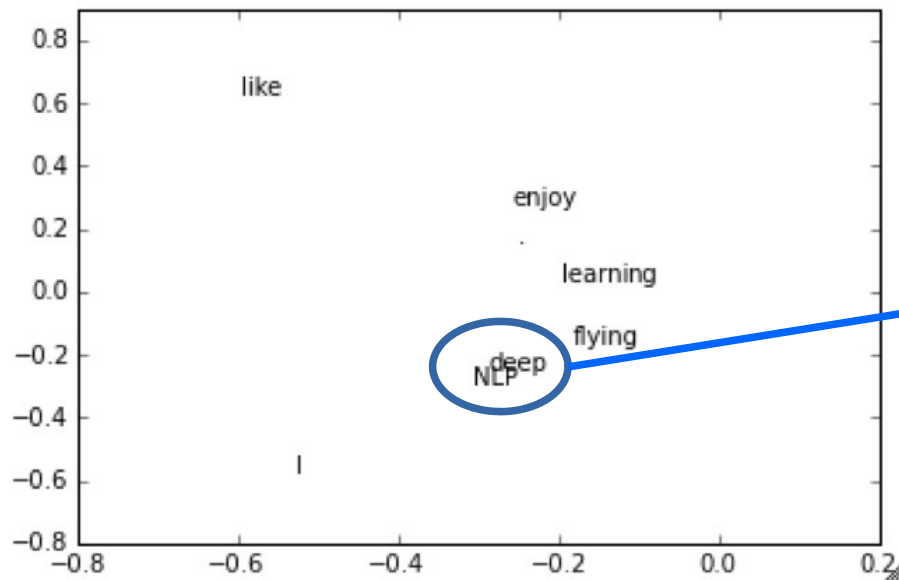
- **Corpus:** I like NLP. I like deep learning. I enjoy flying.



```
for i in xrange(len(words)):
    plt.text(U[i,0], U[i,1], words[i])
```


PLOTTING SVD EXAMPLE IN PYTHON

- **Corpus:** I like NLP. I like deep learning. I enjoy flying.

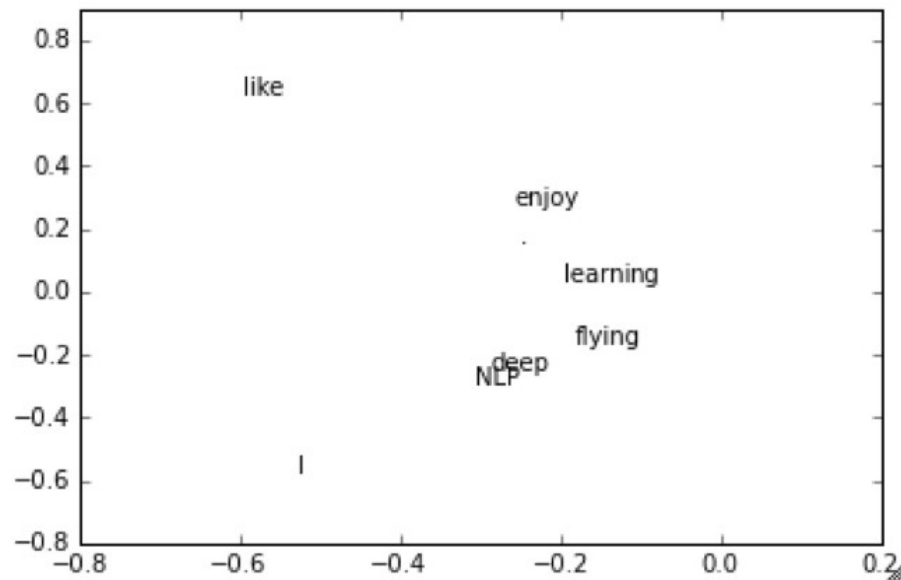


```
for i in xrange(len(words)):
    plt.text(U[i,0], U[i,1], words[i])
```

NLP and deep aren't directly connected in the corpus (Δ is 1), but have common context (like)

COMPUTING WORD SIMILARITY

- **Corpus:** I like NLP. I like deep learning. I enjoy flying.



```
for i in xrange(len(words)):
    plt.text(U[i,0], U[i,1], words[i])
```

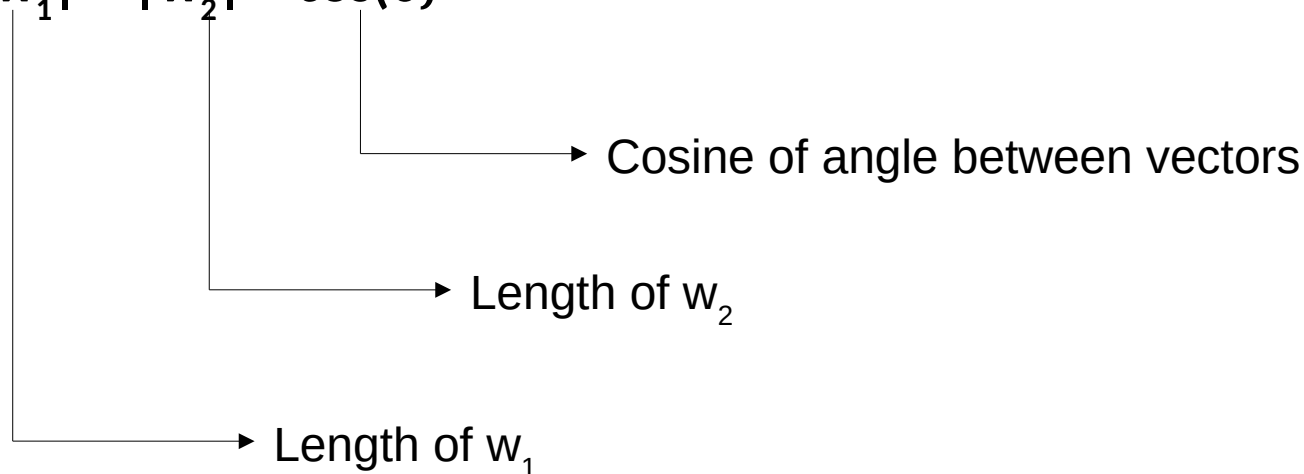
We can compute similarity
between w_i and w_j by comparing:
 $U[i, 0:k]$ and $U[j, 0:k]$

COMPUTING WORD SIMILARITY

- Given 2 words w_1 and w_2 , similarity is computed as:

- Dot/inner product, which equates:

$$|w_1| * |w_2| * \cos(\theta)$$



COMPUTING WORD SIMILARITY

- Given 2 words w_1 and w_2 , similarity is computed as:
 - Dot/inner product, which equates:
 $|w_1| * |w_2| * \cos(\theta)$
 - **High similarity** for:
near-parallel vectors with high values in same dimensions.
 - **Low similarity** for:
orthogonal vectors, low value vectors.

PROS AND CONS OF SVD

- **Pro:** has shown to perform well in a number of tasks.
 - Useful e.g. for topic models, Latent Dirichlet Allocation (LDA).
- **Con:** dimensions need to change as new words are added to the corpus, costly.
- **Con:** resulting vectors can still be high dimensional and sparse.
- **Con:** Quadratic cost to perform SVD.



WORD EMBEDDINGS: STATE-OF-THE-ART ALTERNATIVES TO SVD

ITERATION BASED METHODS: WORD2VEC

- **Main intuition:** Instead of computing co-occurrences from entire corpus, **predict surrounding words** in a window of length c of every word.
 - Allows **easier updates**, faster to incorporate new words in model.
 - Leads to **low dimensional, dense vectors**.
 - This is the idea behind **word2vec** (Mikolov 2013)

ALTERNATIVE: ITERATION BASED METHODS

- **Intuition: predict surrounding words.**

e.g. will you **X** it? → try to predict X, use a neural network to predict and refine predictions.

will you **buy** it? → high score

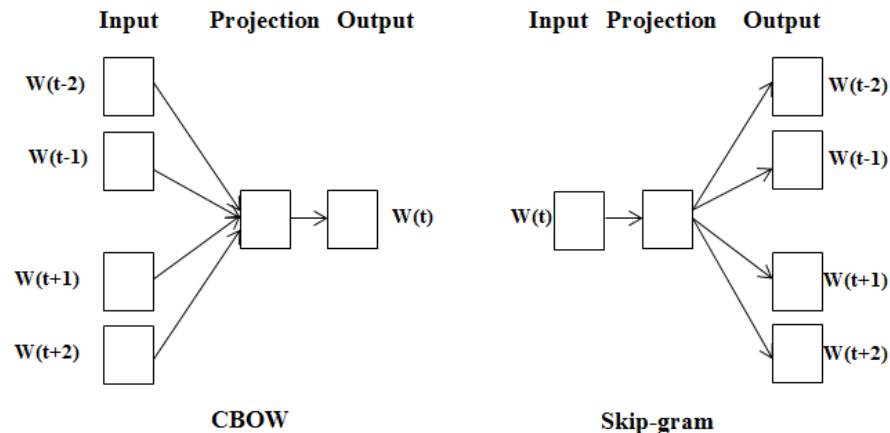
will you **purchase** it? → high score

will you **beer** it? → low score

will you **university** it? → low score

WORD2VEC: CBOW AND SKIPGRAM MODELS

- **Continuous bag of words model (CBOW):** having the context, predict a word.
- **Skip gram model:** having the word, predict its context.



WORD2VEC: WHY IS IT COOL?

- They are very good for encoding similarity.
 - Analogies testing dimensions of similarity can be solved quite well just by doing vector subtraction in the embedding space
- Syntactically
- $x_{apple} - x_{apples} \approx x_{car} - x_{cars} \approx x_{family} - x_{families}$
- Similarly for verb and adjective morphological forms
- Semantically (Semeval 2012 task 2)
- $x_{shirt} - x_{clothing} \approx x_{chair} - x_{furniture}$
 - $x_{king} - x_{man} \approx x_{queen} - x_{woman}$

WORD2VEC: WHY IS IT COOL?

- They are **very good for inferring word relations**:
 - $v(\text{'Paris'}) - v(\text{'France'}) + v(\text{'Italy'}) = v(\text{'Rome'})$
 - $v(\text{'king'}) - v(\text{'man'}) + v(\text{'woman'}) = v(\text{'queen'})$

PROS AND CONS: ITERATION BASED METHODS

- **Pro:** Do not need to operate on entire corpus which involves very sparse matrices.
- **Pro:** Can capture semantic properties of words as linear relationships between word vectors.
- **Pro:** Fast and can be easily updated with new sentences.
- **Con:** Can't take into account the vast amount of repetition in the data.

ANOTHER ALTERNATIVE: GLOVE

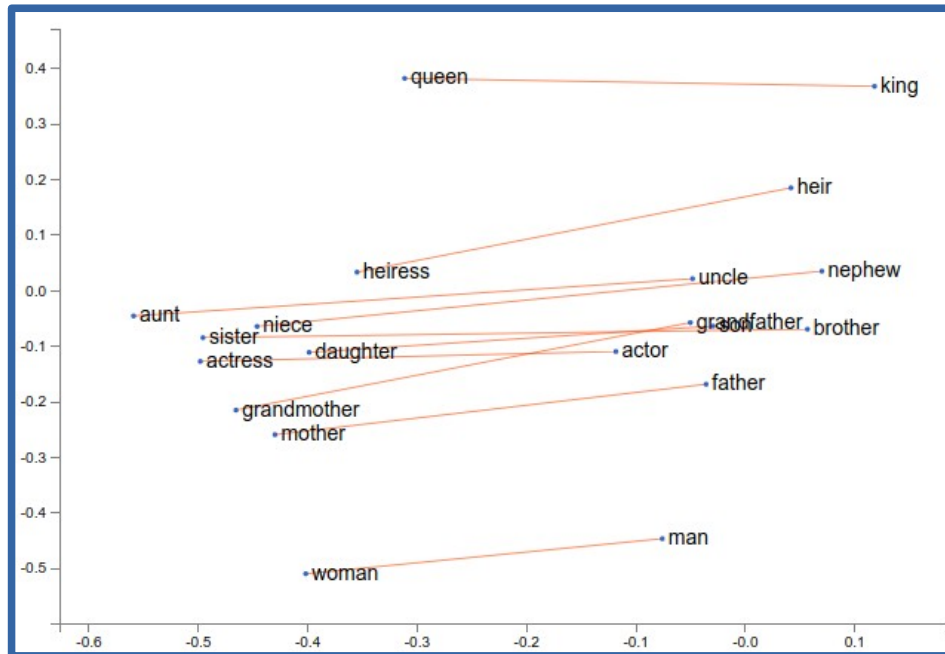
- **Glove (Pennington et al. 2014)**, is similar to word2vec:
 - Count-based method instead of prediction-based.
 - Does matrix factorisation for dimensionality reduction.
- Can leverage repetitions in the corpus as using the entire word co-occurrence matrix.
- **How?** Train only on non-zero entries of the co-occurrence matrix.

GLOVE

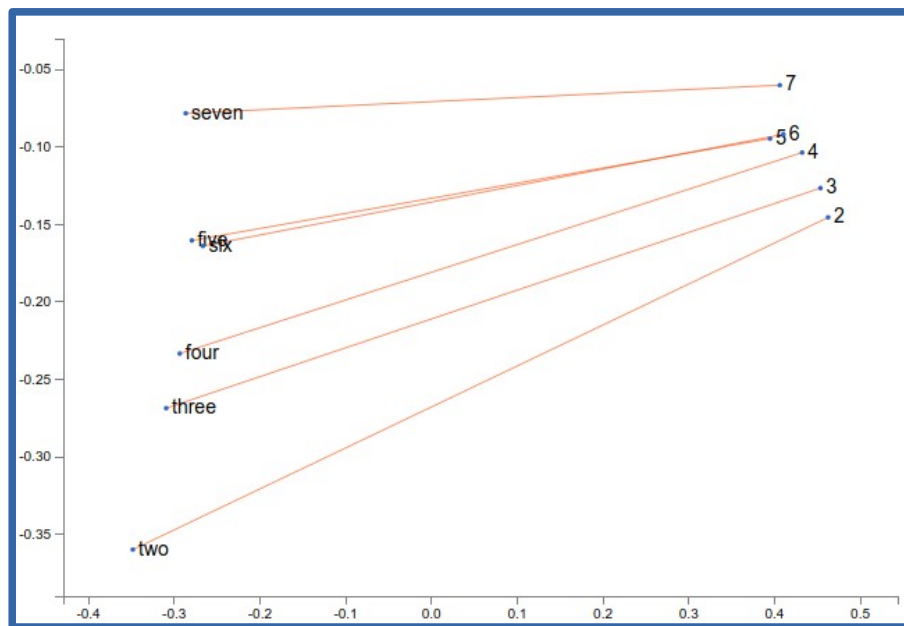
- **Computationally expensive for building matrix, then much faster** as non-zero entries are not so many.
- **Intuition:** relationships between words should be explored in terms of their co-occurrence probabilities with some selected words k .

| Probability and Ratio | $k = \text{solid}$ | $k = \text{gas}$ | $k = \text{water}$ | $k = \text{fashion}$ |
|-------------------------------------|----------------------|----------------------|----------------------|----------------------|
| $P(k \text{ice})$ | 1.9×10^{-4} | 6.6×10^{-5} | 3.0×10^{-3} | 1.7×10^{-5} |
| $P(k \text{steam})$ | 2.2×10^{-5} | 7.8×10^{-4} | 2.2×10^{-3} | 1.8×10^{-5} |
| $P(k \text{ice})/P(k \text{steam})$ | 8.9 | 8.5×10^{-2} | 1.36 | 0.96 |

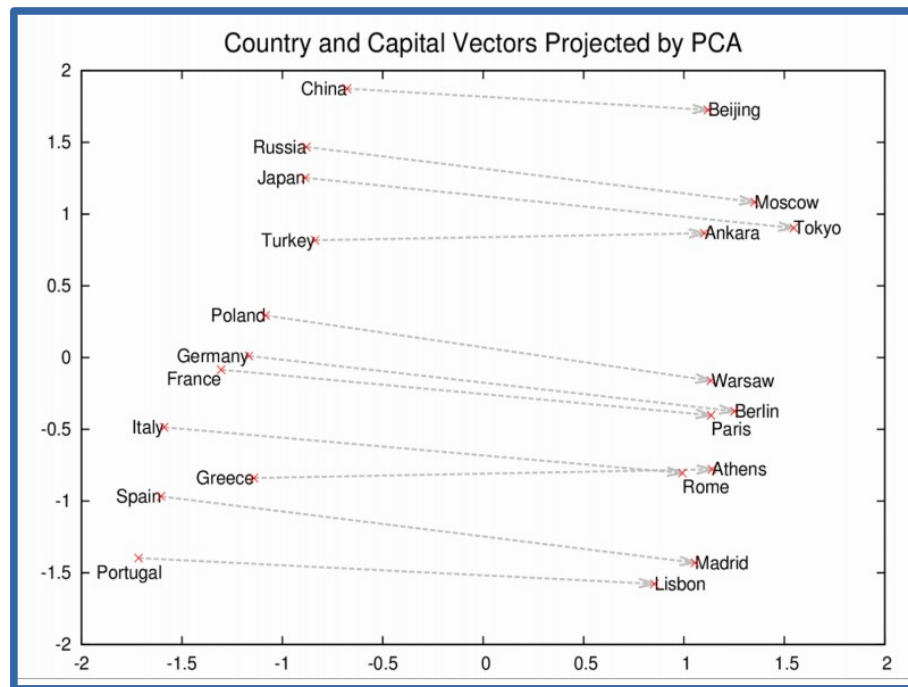
GLOVE: VISUALISATION



GLOVE: VISUALISATION



GLOVE: VISUALISATION



GLOVE: VISUALISATION

- Want to play around?

<https://lamiowce.github.io/word2viz/>

Explore word analogies

What do you want to see?

Gender analogies

Modify words

Type a new word...

Add

Type a new word...

Type a new word...

Add pair

X axis:

young

old

Y axis:

girl

woman

Change axes labels

EVALUATION OF WORD EMBEDDINGS

- **Extrinsic evaluation:** test your model in a text classification, sentiment analysis, machine translation,... task!
 - Does it **outperform other methods** (e.g. bag-of-words)?
 - **Compare two models** A and B: which one's better?

EVALUATION OF WORD EMBEDDINGS

- **Intrinsic evaluation:**

- Use datasets labelled with word similarities:
 - e.g. TOEFL dataset: “pose” is closest in meaning to:
a) claim, b) model, c) assume, d) present

do we get it right with embeddings?

- Common sense:

- Paris + UK – France = London?

PYTHON: USING WORD2VEC

- Preparing the input:
Word2Vec takes lists of lists of words (lists of sentences) as input.

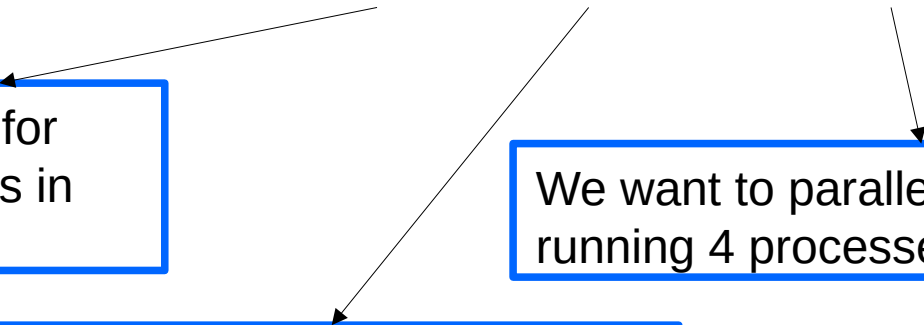
e.g.:

```
sentences = [['this', 'is', 'my', 'first', 'sentence'],  
              ['a', 'short', 'sentence'],  
              ['another', 'sentence'],  
              ['and', 'this', 'is', 'the', 'last', 'one']]
```

PYTHON: USING WORD2VEC

- Training the model:

```
model = Word2Vec(sentences, min_count=10, size=300, workers=4)
```



We will only train vectors for words occurring 10+ times in the corpus

We want to parallelise the task running 4 processes

We want to produce word vectors of 300 dimensions

PYTHON: USING WORD2VEC

- It's memory intensive!

It stores matrices: **#vocabulary** (dependent on min_count), **#size** (size parameter) of **floats** (single precision aka 4 bytes).

Three such matrices are held in RAM. If you have:
100,000 unique words, size=200, the model will require approx.:

$100,000 * 200 * 4 * 3 \text{ bytes} = \sim 229\text{MB}.$

PYTHON: USING WORD2VEC

- Storing a model:

```
model = Word2Vec.load_word2vec_format('mymodel.txt', binary=False)
```

- or

```
model = Word2Vec.load_word2vec_format('mymodel.bin.gz', binary=True)
```

- Resuming training:

```
model = gensim.models.Word2Vec.load('mymodel.bin.gz')  
model.train(more_sentences)
```


PYTHON: USING WORD2VEC

Using the model

Word2vec supports several word similarity tasks out of the box:

```
1 | model.most_similar(positive=['woman', 'king'], negative=['man'], topn=1)
2 | [('queen', 0.50882536)]
3 | model.doesnt_match("breakfast cereal dinner lunch".split())
4 | 'cereal'
5 | model.similarity('woman', 'man')
6 | 0.73723527
```

If you need the raw output vectors in your application, you can access these either on a word-by-word basis

```
1 | model['computer'] # raw NumPy vector of a word
2 | array([-0.00449447, -0.00310097,  0.02421786, ...], dtype=float32)
```

PYTHON: USING WORD2VEC

```
1 | model['computer'] # raw NumPy vector of a word
2 | array([-0.00449447, -0.00310097,  0.02421786, ...], dtype=float32)
```

- This will give us the vector representation of 'computer:'
 $v(\text{'computer'}) = \{-0.00449447, -0.00310097, \dots\}$
- How do we get then the vector representations for sentences, e.g.:
I have installed Ubuntu on my computer

PYTHON: USING WORD2VEC

- **Vector representations for sentences**, e.g.:
 - I have installed Ubuntu on my computer
- Standard practice is either of:
 - **Summing word vectors** (they have the same dimensionality):
 $v('I') + v('have') + v('installed') + v('Ubuntu') + \dots$
 - **Getting the average of word vectors**:
 $(v('I') + v('have') + v('installed') + \dots) / 7$

PRE-TRAINED WORD VECTORS

- One can train a model from a large corpus (millions, if not billions of sentences). Can be time-consuming, memory-intensive.
- **Pre-trained models are available.**
 - **Remember to choose a suitable pre-trained model.**
 - Don't use word vectors pre-trained from news articles when you're working with social media!

PRE-TRAINED WORD VECTORS

- Glove's pre-trained vectors:
<https://nlp.stanford.edu/projects/glove/>

Download pre-trained word vectors

- Pre-trained word vectors. This data is made available under the [Public Domain Dedication and License](http://www.opendatacommons.org/licenses/pddl/1.0/) v1.0 whose <http://www.opendatacommons.org/licenses/pddl/1.0/>.
 - [Wikipedia 2014 + Gigaword 5](#) (6B tokens, 400K vocab, uncased, 50d, 100d, 200d, & 300d vectors, 822 MB download): [glove.42B.300d.zip](#)
 - Common Crawl (42B tokens, 1.9M vocab, uncased, 300d vectors, 1.75 GB download): [glove.42B.300d.zip](#)
 - Common Crawl (840B tokens, 2.2M vocab, cased, 300d vectors, 2.03 GB download): [glove.840B.300d.zip](#)
 - Twitter (2B tweets, 27B tokens, 1.2M vocab, uncased, 25d, 50d, 100d, & 200d vectors, 1.42 GB download): [glove.twitter.27B.zip](#)
- Ruby [script](#) for preprocessing Twitter data

PRE-TRAINED WORD VECTORS

- Pre-trained word vectors for 30+ languages (from Wikipedia):
<https://github.com/Kyubyong/wordvectors>

| Language | ISO 639-1 | Vector Size | Corpus Size | Vocabulary Size |
|---------------------------------|-----------|-------------|-------------|-----------------|
| Bengali (w) Bengali (f) | bn | 300 | 147M | 10059 |
| Catalan (w) Catalan (f) | ca | 300 | 967M | 50013 |
| Chinese (w) Chinese (f) | zh | 300 | 1G | 50101 |
| Danish (w) Danish (f) | da | 300 | 295M | 30134 |
| Dutch (w) Dutch (f) | nl | 300 | 1G | 50160 |
| Esperanto (w) Esperanto (f) | eo | 300 | 1G | 50597 |
| Finnish (w) Finnish (f) | fi | 300 | 467M | 30029 |
| French (w) French (f) | fr | 300 | 1G | 50130 |
| German (w) German (f) | de | 300 | 1G | 50006 |
| Hindi (w) Hindi (f) | hi | 300 | 323M | 30393 |
| Hungarian (w) Hungarian (f) | hu | 300 | 692M | 40122 |
| Indonesian (w) Indonesian (f) | id | 300 | 402M | 30048 |

PRE-TRAINED WORD VECTORS

- UK Twitter word embeddings:

https://figshare.com/articles/UK_Twitter_word_embeddings_II_/5791650

UK Twitter word embeddings (II)

16.01.2018, 23:08 by *Vasileios Lamos*

Word embeddings trained on UK Twitter content (II)

The total number of tweets used was approximately 1.1 billion, covering the years 2012 to and including 2016.

Settings: Skip-gram with negative sampling (10 noise words), a window of 9 words, 512 layers (dimensionality) and 10 epochs of training.

PRE-TRAINED WORD VECTORS

- Twitter Word2Vec model:
<https://www.fredericgodin.com/software/>

Twitter Word2vec model

As part of our ACL W-NUT 2015 shared task paper, we release a Twitter word2vec model trained on 400 million tweets, as described in detail in this [paper](#). The model, including Python code to load and access it, can be downloaded [here](#).

REFERENCES

- Gensim (word2vec):
<https://radimrehurek.com/gensim/>
- Word2vec tutorial:
<https://rare-technologies.com/word2vec-tutorial/>
- FastText:
<https://github.com/facebookresearch/fastText/>
- GloVe: Global Vectors for Word Representation:
<https://nlp.stanford.edu/projects/glove/>

ASSOCIATED READING

- Jurafsky, Daniel, and James H. Martin. 2009. Speech and Language Processing: An Introduction to Natural Language Processing, Speech Recognition, and Computational Linguistics. 3rd edition. **Chapter 6.3-6.13.**