

CS918: LECTURE 12

Introduction to Information Retrieval

Arkaitz Zubiaga, 12th November, 2018

LECTURE 12: CONTENTS

- What is Information Retrieval (IR)?
- Indexing Documents.
- Query Processing.
- Positional Indices.
- Other Challenges in Information Retrieval.

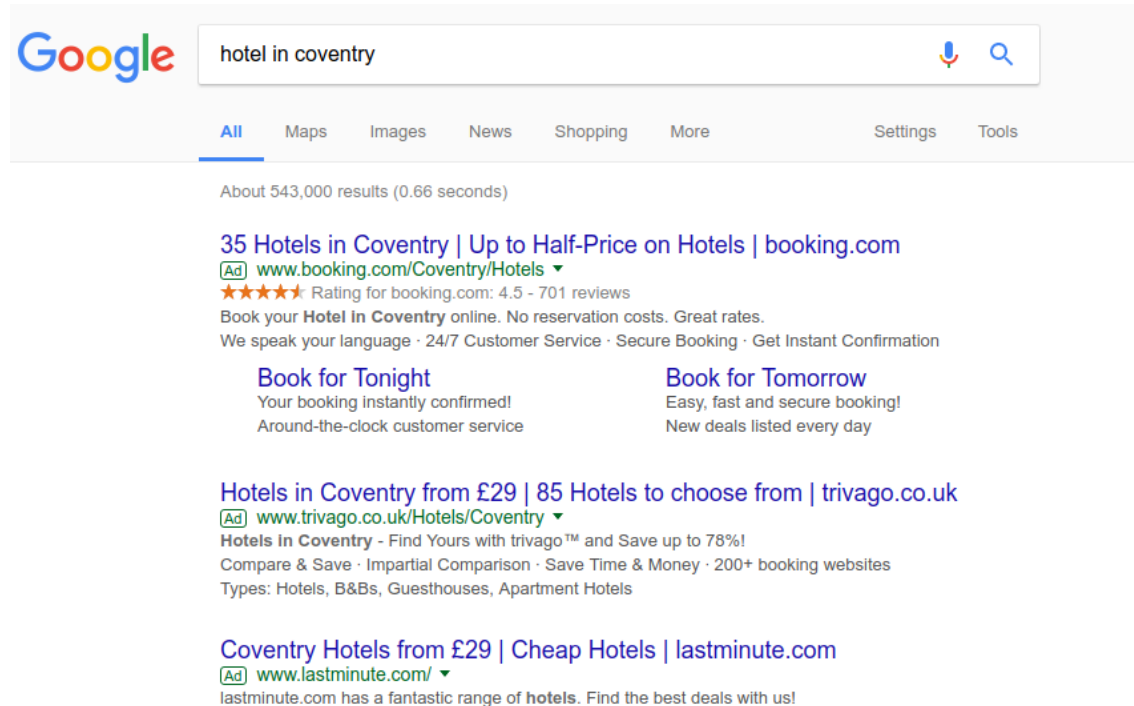
INFORMATION RETRIEVAL

- **Information Retrieval (IR):** from a large collection, the task of obtaining documents that satisfy an **information need**.
 - Collection (e.g. the Web) may include images and videos.
 - In this module we'll **focus on text**.

EXAMPLES OF INFORMATION RETRIEVAL

- **Web search**, e.g. Google.
- **Vertical search**, web search on a particular topic, e.g. Yelp.
- **Email search**.
- Searching content in large **databases or hard drives**.

EXAMPLE OF AN IR SYSTEM?



Google

hotel in coventry

All Maps Images News Shopping More Settings Tools

About 543,000 results (0.66 seconds)

35 Hotels in Coventry | Up to Half-Price on Hotels | booking.com

Ad www.booking.com/Coventry/Hotels ▼

★★★★★ Rating for booking.com: 4.5 - 701 reviews

Book your **Hotel in Coventry** online. No reservation costs. Great rates.

We speak your language · 24/7 Customer Service · Secure Booking · Get Instant Confirmation

<p>Book for Tonight</p> <p>Your booking instantly confirmed!</p> <p>Around-the-clock customer service</p>	<p>Book for Tomorrow</p> <p>Easy, fast and secure booking!</p> <p>New deals listed every day</p>
--	---

Hotels in Coventry from £29 | 85 Hotels to choose from | trivago.co.uk

Ad www.trivago.co.uk/Hotels/Coventry ▼

Hotels In Coventry - Find Yours with trivago™ and Save up to 78%!

Compare & Save · Impartial Comparison · Save Time & Money · 200+ booking websites

Types: Hotels, B&Bs, Guesthouses, Apartment Hotels

Coventry Hotels from £29 | Cheap Hotels | lastminute.com

Ad www.lastminute.com/ ▼

lastminute.com has a fantastic range of hotels. Find the best deals with us!

INFORMATION NEED

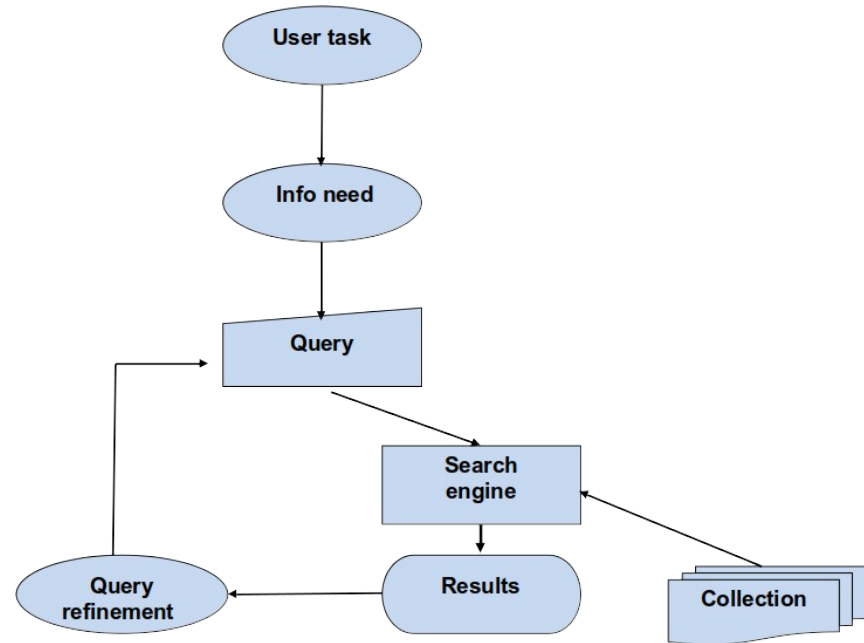
- The **information need** that an IR system has to satisfy is usually expressed as a (short) text **query**, e.g. hotel coventry.
- Many **queries are vague**, i.e. the average search query has 2-4 keywords (Arantzapis & Kamps, 2008).

TYPES OF INFORMATION NEED

- Queries can seek **3 types of information need** (Broder, 2002):
 - **Navigational** (looking for website, e.g. Facebook).
 - **Informational** (looking for info, e.g. Thai food).
 - **Transactional** (buying/looking for product/service, e.g. iPhone 10).
- Google redefined them as: **do, know, go**.
- **Vast majority** (up to 80%) tend to be **informational** or **know**.

THE CLASSIC SEARCH MODEL

- The **search** process **can be iterative**.
 - **Query** gives **results**.
 - If **unhappy** with results, **refine the query**.
- With a good IR system, we aim to **minimise the number of times the user refines the query**.



RELEVANCE TO A QUERY

- In IR, documents in a collection are deemed **relevant (R)** or **non-relevant (N)** with **respect to a particular query**.
 - Or sometimes levels of relevance, e.g. 1-5.
- The objective of an IR system is to present, for a given query, **as many relevant results as possible** (and as few non-relevant results as possible).
 - **58%** of users **never go to the 2nd page** of search results (Jansen et al., 1998)

SEARCHING THROUGH LARGE COLLECTIONS

- In a **small collection**, we can **process files on the fly**.
 - For a given query, go through all files and check if the query text appears in each file.
 - It would **take ages for large collections**.
- For **large collections**, **indexing** is the solution.
 - i.e. pregenerating lists of word-document associations.



WARWICK

INDEXING DOCUMENTS

TERM-DOCUMENT INCIDENCE MATRICES

- Reduced sample of Shakespeare's works:

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0

TERM-DOCUMENT INCIDENCE MATRICES

- Reduced sample of Shakespeare's works:

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0


- We can search for: +Brutus +Caesar -Calpurnia (2 results)

EVEN BIGGER COLLECTIONS

- Collection of **$N=1\text{M}$ documents, each with 1K words.**
 - Say there are **$M = 500\text{K}$ distinct terms** among these.
- $500\text{K} \times 1\text{M}$ matrix has **half-a-trillion 0's and 1's.**
 - **Very sparse**, only one billion 1's (that's 0.2% of the values).
 - Alternative: **record only 1's** \rightarrow use an **inverted index**.

INVERTED INDEX

- For each term t , **store list of documents that contain t** .
 - List needs to have **variable size**.



- Word 'Brutus' occurs in documents ID 1, ID 2, ID 4, etc.

INDEXING DOCUMENTS: EXAMPLE

- List (token, Document ID) pairs.

Doc 1

I did enact Julius
Caesar I was killed
i' the Capitol;
Brutus killed me.

Doc 2

So let it be with
Caesar. The noble
Brutus hath told you
Caesar was ambitious



Term	docID
I	1
did	1
enact	1
julius	1
caesar	1
I	1
was	1
killed	1
i'	1
the	1
capitol	1
brutus	1
killed	1
me	1
so	2
let	2
it	2
be	2
with	2
caesar	2
the	2
noble	2
brutus	2
hath	2
told	2
you	2
caesar	2
was	2
ambitious	2

INDEXING DOCUMENTS: EXAMPLE

- Sort (alphabetically) by tokens and then doc ID.

Term	docID		Term	docID
I	1		ambitious	2
did	1		be	2
enact	1		brutus	1
julius	1		brutus	2
caesar	1		capitol	1
I	1		caesar	1
was	1		caesar	2
killed	1		caesar	2
i'	1		did	1
the	1		enact	1
capitol	1		hath	1
brutus	1		I	1
killed	1		I	1
me	1		i'	1
so	2		it	2
let	2		julius	1
it	2		killed	1
be	2		killed	1
with	2		let	2
caesar	2		me	1
the	2		noble	2
noble	2		so	2
brutus	2		the	1
hath	2		the	2
told	2		told	2
you	2		you	2
caesar	2		was	1
was	2		was	2
ambitious	2		with	2

INDEXING DOCUMENTS: EXAMPLE

- Merge entries + add frequency counts: **dictionary** and **postings**.

Term	docID
ambitious	2
be	2
brutus	1
brutus	2
capitol	1
caesar	1
caesar	2
caesar	2
did	1
enact	1
hath	1
I	1
I	1
i'	1
it	2
julius	1
killed	1
killed	1
let	2
me	1
noble	2
so	2
the	1
the	2
told	2
you	2
was	1
was	2
with	2



term	doc. freq.	→	postings lists
ambitious	1	→	2
be	1	→	2
brutus	2	→	1 → 2
capitol	1	→	1
caesar	2	→	1 → 2
did	1	→	1
enact	1	→	1
hath	1	→	2
i	1	→	1
i'	1	→	1
it	1	→	2
julius	1	→	1
killed	1	→	1
let	1	→	2
me	1	→	1
noble	1	→	2
so	1	→	2
the	2	→	1 → 2
told	1	→	2
you	1	→	2
was	2	→	1 → 2
with	1	→	2




WARWICK

QUERY PROCESSING

QUERY PROCESSING: AND

- Search query: **Brutus AND Caesar**
 - Retrieve postings with **Brutus**.
 - Retrieve postings with **Caesar**.
 - Get the **intersection** as the set of results (docs 2 and 8).



GETTING THE INTERSECTION

```
INTERSECT( $p_1, p_2$ )  
1   $answer \leftarrow \langle \rangle$   
2  while  $p_1 \neq \text{NIL}$  and  $p_2 \neq \text{NIL}$   
3  do if  $\text{docID}(p_1) = \text{docID}(p_2)$   
4      then  $\text{ADD}(answer, \text{docID}(p_1))$   
5           $p_1 \leftarrow \text{next}(p_1)$   
6           $p_2 \leftarrow \text{next}(p_2)$   
7      else if  $\text{docID}(p_1) < \text{docID}(p_2)$   
8          then  $p_1 \leftarrow \text{next}(p_1)$   
9          else  $p_2 \leftarrow \text{next}(p_2)$   
10 return  $answer$ 
```

GETTING THE INTERSECTION

```
INTERSECT( $p_1, p_2$ )
1   $answer \leftarrow \langle \rangle$ 
2  while  $p_1 \neq \text{NIL}$  and  $p_2 \neq \text{NIL}$ 
3  do if  $\text{docID}(p_1) = \text{docID}(p_2)$ 
4      then  $\text{ADD}(answer, \text{docID}(p_1))$ 
5           $p_1 \leftarrow \text{next}(p_1)$ 
6           $p_2 \leftarrow \text{next}(p_2)$ 
7      else if  $\text{docID}(p_1) < \text{docID}(p_2)$ 
8          then  $p_1 \leftarrow \text{next}(p_1)$ 
9          else  $p_2 \leftarrow \text{next}(p_2)$ 
10 return  $answer$ 
```

- Complexity of algorithm: $O(m+n)$, as we iterate through all items – m is the length of p_1 , and n is the length of p_2

SKIP POINTERS

- **Can we** improve the linear time $O(m+n)$ and **compute it in sublinear time?**
 - We can use **skip pointers**, i.e.:
 - If p_1 has: 1, 3, 5, 15, 40
 - and p_2 has: 35, 40, 90
 - Pointers will initially be $pp_1 = 1$ and $pp_2 = 35$
 - We can skip all values smaller than 35 in p_1 .

SKIP POINTERS

```
INTERSECTWITHSKIPS( $p_1, p_2$ )
1   $answer \leftarrow \langle \rangle$ 
2  while  $p_1 \neq \text{NIL}$  and  $p_2 \neq \text{NIL}$ 
3  do if  $\text{docID}(p_1) = \text{docID}(p_2)$ 
4      then  $\text{ADD}(answer, \text{docID}(p_1))$ 
5           $p_1 \leftarrow \text{next}(p_1)$ 
6           $p_2 \leftarrow \text{next}(p_2)$ 
7  else if  $\text{docID}(p_1) < \text{docID}(p_2)$ 
8      then if  $\text{hasSkip}(p_1)$  and  $(\text{docID}(\text{skip}(p_1)) \leq \text{docID}(p_2))$ 
9          then while  $\text{hasSkip}(p_1)$  and  $(\text{docID}(\text{skip}(p_1)) \leq \text{docID}(p_2))$ 
10             do  $p_1 \leftarrow \text{skip}(p_1)$ 
11             else  $p_1 \leftarrow \text{next}(p_1)$ 
12         else if  $\text{hasSkip}(p_2)$  and  $(\text{docID}(\text{skip}(p_2)) \leq \text{docID}(p_1))$ 
13             then while  $\text{hasSkip}(p_2)$  and  $(\text{docID}(\text{skip}(p_2)) \leq \text{docID}(p_1))$ 
14                 do  $p_2 \leftarrow \text{skip}(p_2)$ 
15                 else  $p_2 \leftarrow \text{next}(p_2)$ 
16 return  $answer$ 
```


PHRASE QUERIES

- I want to search for “**University of Warwick**” – as a phrase.
- For this query, “You can live in Warwick if you are a student at the university” is **NOT a match**.
- If we want to do this search, then our **<term : docs> index is not enough**.

HOW ABOUT BIGRAM INDICES?

- **Index bigrams** instead of single words.
- For the document, “I went to the University of Warwick”, index:
I went, went to, to the, the University, etc.
- We can **easily search for bigrams now**, but we **can’t look for “University of Warwick”** yet!
 - We can search for “University of AND of Warwick”, but there is no guarantee that they occur together in the document.

LONGER N-GRAM INDICES?

- **Longer n-gram indices are not feasible**, too many possible n-grams.
- **Bigrams** could be used when we don't need longer search queries, but it's **generally not enough**.
- Use of **bigrams to look for longer n-grams can lead to false positives** (as with the “University of AND of Warwick” example)



WARWICK

POSITIONAL INDICES

THE SOLUTION: POSITIONAL INDICES

- In the postings, **store for each term the position(s)** in which tokens of it appear:

<term, number of docs containing term;

doc1: position1, position2 ... ;

doc2: position1, position2 ... ;

etc.>

POSITIONAL INDEX: EXAMPLE

I have two documents:

- doc1: University of Warwick
 - doc2: Warwick University
-
- My index:
 - <University, 2; doc1: 1, doc2: 2;
 - Warwick, 2; doc1: 3, doc2: 1;
 - of, 1; doc1: 2>

POSITIONAL INDEX: EXAMPLE

- The search query could be: **“to be or not to be”**
 - Which documents contain it?
- For phrase queries, we **use a merge algorithm recursively at the document level.**
- But we now **need to deal with more than just equality.**

POSITIONAL INDEX: EXAMPLE

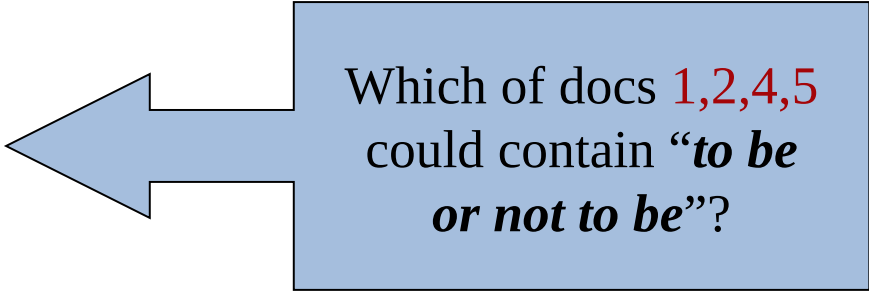
<**be**: 993427;

1: 7, 18, 33, 72, 86, 231;

2: 3, 149;

4: 17, 191, 291, 430, 434;

5: 363, 367, ...>



Which of docs **1,2,4,5**
could contain “*to be*
or not to be”?

POSITIONAL INDEX: EXAMPLE

- Get **inverted index entries for each distinct term**: to, be, or, not.
- Merge their doc:position lists.
 - to:
2:1,17,74,222,551; 4:8,16,190,**429,433**; 7:13,23,191; ...
 - be:
1:17,19; 4:17,191,291,**430,434**; 5:14,19,101; ...
 - or...
- Look for **occurrences where the positions match the sequence** of our query: “to be or not to be”.

POSITIONAL INDEX: EXAMPLE

- Get **inverted index entries for each distinct term**: to, be, or, not.
- Merge their doc:position lists.

429: to, 430: be, 431: or, 432: not, 433: to, 434: be

- to:

2:1,17,74,222,551; 4:8,16,190,**429,433**; 7:13,23,191; ...

- be:

1:17,19; 4:17,191,291,**430,434**; 5:14,19,101; ...

- or...

- Look for **occurrences where the positions match the sequence** of our query: “to be or not to be”.

BEYOND A DOCUMENT'S CONTENT

- What if a document doesn't have the query keywords but it is relevant?

e.g. if I search for “University in the West Midlands”

warwick.ac.uk may not contain those words.

BEYOND A DOCUMENT'S CONTENT

- What if a document doesn't have the query keywords but it is relevant?

Web search engines use “anchor texts” from incoming web links.

Intuition:

Warwick itself is unlikely to say “University in the West Midlands”

Another website may say “check out this [university in the West Midlands](#)”, with a link to Warwick.

BEYOND A DOCUMENT'S CONTENT

- OK, but what happens if an anchor text linking to Warwick says “[click here](#)”.
 - Our IR system will believe Warwick is a relevant result for the “click here” search query.
- We will need to determine the **keywords** that are **meaningful**, by **weighting** them.
 - Forthcoming lectures.

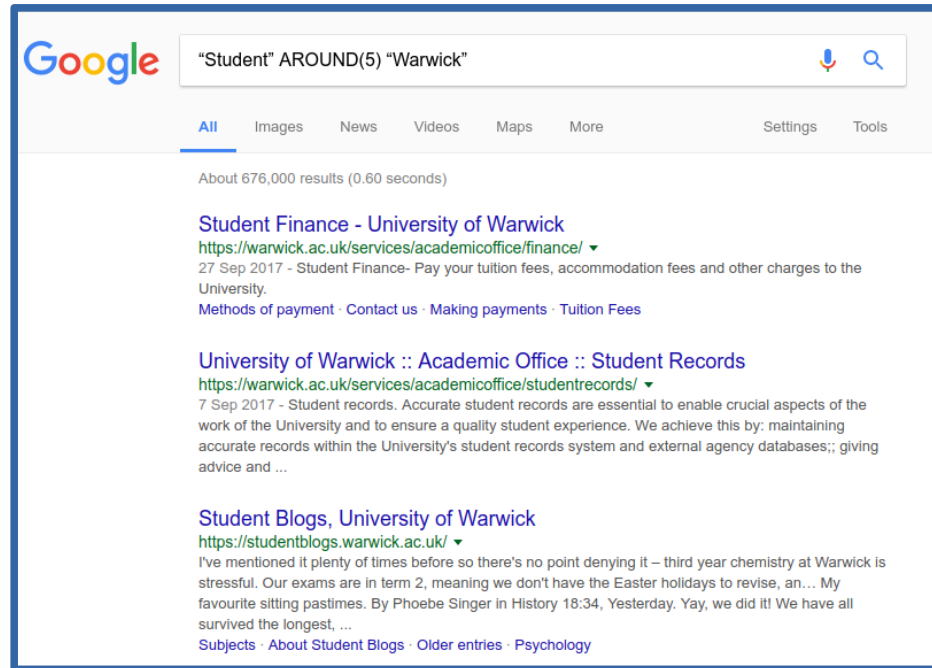


OTHER CHALLENGES IN INFORMATION RETRIEVAL

PROXIMITY QUERIES

- We may still want to do **more sophisticated queries, e.g. proximity queries.**
 - Search for two phrases, which occur within a maximum distance of K words between them.

PROXIMITY QUERIES



PROXIMITY QUERIES

- Query: **Student AROUND(5) Warwick.**
 - We can do this with positional indices.
 - With bigram indices we can't.
- The **algorithm for getting the intersection** here is more complex:
 - We need to get the intersection of “student” and “Warwick”.
 - with the restriction that there has to be a maximum of 5 words in between.

PROXIMITY QUERIES: ALGORITHM

```

POSITIONALINTERSECT( $p_1, p_2, k$ )
1   $answer \leftarrow \langle \rangle$ 
2  while  $p_1 \neq \text{NIL}$  and  $p_2 \neq \text{NIL}$ 
3  do if  $\text{docID}(p_1) = \text{docID}(p_2)$ 
4      then  $l \leftarrow \langle \rangle$ 
5           $pp_1 \leftarrow \text{positions}(p_1)$ 
6           $pp_2 \leftarrow \text{positions}(p_2)$ 
7          while  $pp_1 \neq \text{NIL}$ 
8          do while  $pp_2 \neq \text{NIL}$ 
9              do if  $|\text{pos}(pp_1) - \text{pos}(pp_2)| \leq k$ 
10                  then  $\text{ADD}(l, \text{pos}(pp_2))$ 
11                  else if  $\text{pos}(pp_2) > \text{pos}(pp_1)$ 
12                      then break
13                       $pp_2 \leftarrow \text{next}(pp_2)$ 
14                  while  $l \neq \langle \rangle$  and  $|l[0] - \text{pos}(pp_1)| > k$ 
15                      do  $\text{DELETE}(l[0])$ 
16                      for each  $ps \in l$ 
17                          do  $\text{ADD}(answer, \langle \text{docID}(p_1), \text{pos}(pp_1), ps \rangle)$ 
18                       $pp_1 \leftarrow \text{next}(pp_1)$ 
19                   $p_1 \leftarrow \text{next}(p_1)$ 
20                   $p_2 \leftarrow \text{next}(p_2)$ 
21          else if  $\text{docID}(p_1) < \text{docID}(p_2)$ 
22              then  $p_1 \leftarrow \text{next}(p_1)$ 
23              else  $p_2 \leftarrow \text{next}(p_2)$ 
24  return  $answer$ 
  
```

WILDCARD SEARCHES

- Query: **University * Warwick**.
- We want to look for **just one word** in between (the meaning of * is different from regex here).
- Positional indices can handle this, just as with **proximity queries where K=1**.

POSITIONAL INDEX SIZE

- Positional index uses **more space** than a binary index.
- However, positional indices are today's **standard approach to index documents**, given their **flexibility** for searching.

POSITIONAL INDEX SIZE

- Rules of thumb:
 - A **positional index** is **2-4 as large** as a non-positional index.
 - Positional index size **35-50% of volume of original text**.
 - **Caveat:** all of this holds for “English-like” languages.

COMBINATION SCHEMES

- Positional indices and bigram indices **can be combined** to get the most of each approach:
 - For **frequent phrases** (“Michael Jackson”, “Britney Spears”, “The Who”) it is inefficient to keep on merging positional postings lists.
- For **very popular search queries**, you can also **cache** the results.

COMBINATION SCHEMES

- Williams et al. (2004) evaluated a **mixed indexing scheme**.
 - A typical web query mixture was **executed in $\frac{1}{4}$ of the time** of using just a positional index.
 - It required **26% more space** than having a positional index alone.

MORE POSSIBLE SEARCHES

- **Case sensitive search:** if we lowercase everything before indexing, we can't consider the case when searching.
- **Search page titles only:** need an additional flag to indicate that word position belongs to the title.
- **Search numeric ranges:** e.g. £50..£100. If numbers are just another string in our indices, we won't be able to search for this.

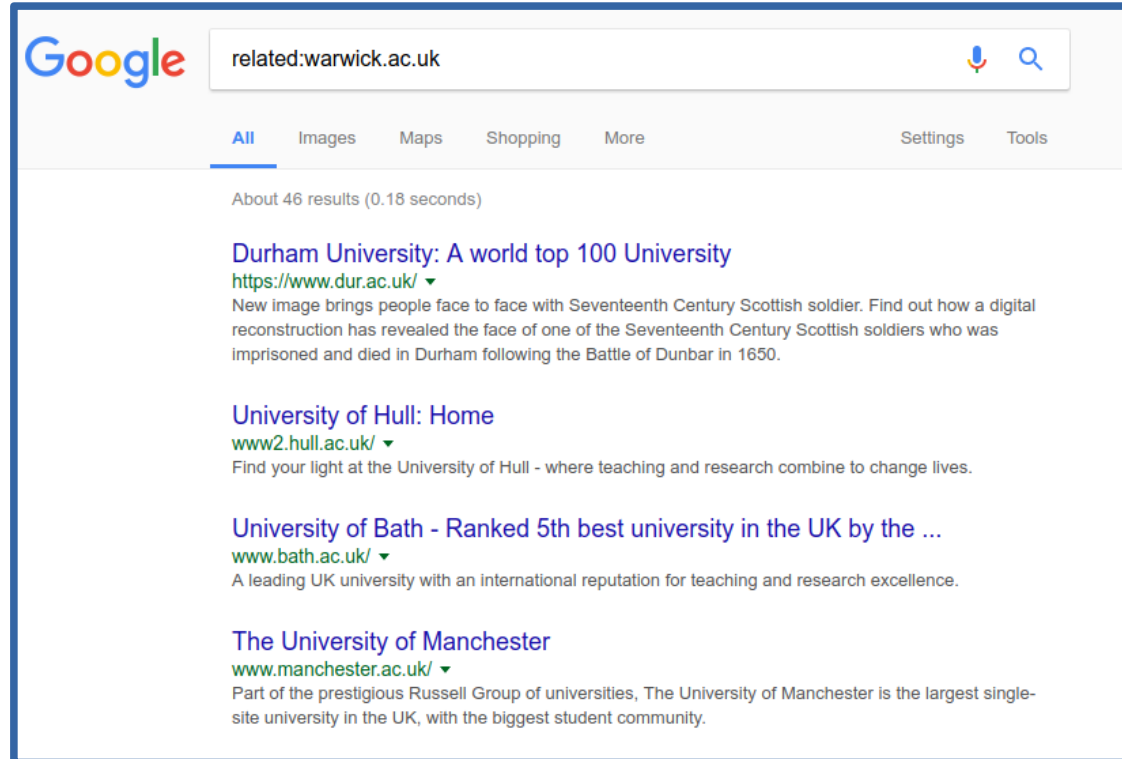
MORE CHALLENGES

- For all these searches, the **challenge lies in implementing it efficiently**:
 - **Case: we could store original and lowercased words**, but that **increases the index size** substantially.
 - We could also have **separate indices for page titles and for numeric values**, but again, that's much more data to store.

MORE CHALLENGES

- With **Google**, you can even search for **related pages**:
e.g. `related:warwick.ac.uk`

MORE CHALLENGES



MORE CHALLENGES

- With **Google**, you can even search for **related pages**:
e.g. `related:warwick.ac.uk`
returns a bunch of university websites
- They may also be **storing content similarity between websites**.
 - Again, that's **much more data for the index**.

MORE CHALLENGES

- So far, we **can retrieve documents that match a query.**
- That's fine, but we often get **many results.**
 - And **we want to rank them.**
 - **Weigh them based on relevance**, not just binary match.
- Ranking is an additional challenge in information retrieval.
 - Next lecture!

RESOURCES

- Apache Lucene (open source search software, Java):
<https://lucene.apache.org/>
- Python wrapper for Lucene:
<http://lucene.apache.org/pylucene/>
- Apache Nutch (web crawler, can be integrated with Lucene to build a search engine):
<http://nutch.apache.org/>

ASSOCIATED READING

- Manning, C. D., Raghavan, P., & Schütze, H. (2008). Introduction to information retrieval (Vol. 1, No. 1, p. 496). Cambridge: Cambridge university press. **Chapters 1-2.**
<https://nlp.stanford.edu/IR-book/pdf/irbookonlinereading.pdf>